

# PyTorch under the hood

## A guide to understand PyTorch internals



Christian S. Perone  
(*christian.perone@gmail.com*)  
<http://blog.christianperone.com>

PyData Montreal, Feb 2019

# Agenda

## TENSORS

- Tensors

- Python objects

- Zero-copy

- Tensor storage

- Memory allocators (CPU/GPU)

- The big picture

## JIT

- Just-in-time compiler

- Tracing

- Scripting

- Why TorchScript ?

- Building IR and JIT Phases

- Optimizations

- Serialization

- Using models in other languages

## PRODUCTION

- Some tips

## Q&A

# WHO AM I

- ▶ **Christian S. Perone**
- ▶ 14 years working with Machine Learning, Data Science and Software Engineering in industry R&D
- ▶ Blog at
  - ▶ [blog.christianperone.com](http://blog.christianperone.com)
  - ▶ Open-source projects at
    - ▶ <https://github.com/perone>
    - ▶ Twitter @tarantulae



# DISCLAIMER

- PyTorch is a **moving target**, Deep Learning ecosystem moves fast and big changes happens every week;

# DISCLAIMER

- ▶ PyTorch is a **moving target**, Deep Learning ecosystem moves fast and big changes happens every week;
- ▶ This is not a talk to teach you the basics of PyTorch or how to train your network, but to teach you how **PyTorch components works under the hood** in a intuitive way;

# DISCLAIMER

- ▶ PyTorch is a **moving target**, Deep Learning ecosystem moves fast and big changes happens every week;
- ▶ This is not a talk to teach you the basics of PyTorch or how to train your network, but to teach you how **PyTorch components works under the hood** in a intuitive way;
- ▶ This talk is updated to the PyTorch v.1.0.1 version;

## Section I

## ∞ TENSORS ∞

# TENSORS

Simply put, TENSORS are a generalization of vectors and matrices. In PyTorch, they are a multi-dimensional matrix containing elements of a **single data type**.



# TENSORS

Simply put, TENSORS are a generalization of vectors and matrices. In PyTorch, they are a multi-dimensional matrix containing elements of a **single data type**.

```
>>> import torch
>>> t = torch.tensor([[1., -1.], [1., -1.]])
>>> t
tensor([[ 1., -1.]
         [ 1., -1.]])
```

# TENSORS

Simply put, TENSORS are a generalization of vectors and matrices. In PyTorch, they are a multi-dimensional matrix containing elements of a **single data type**.

```
>>> import torch
>>> t = torch.tensor([[1., -1.], [1., -1.]])
>>> t
tensor([[ 1., -1.]
         [ 1., -1.]])

>>> t.dtype # They have a type
torch.float32
```

# TENSORS

Simply put, TENSORS are a generalization of vectors and matrices. In PyTorch, they are a multi-dimensional matrix containing elements of a **single data type**.

```
>>> import torch
>>> t = torch.tensor([[1., -1.], [1., -1.]])
>>> t
tensor([[ 1., -1.]
         [ 1., -1.]])

>>> t.dtype # They have a type
torch.float32

>>> t.shape # a shape
torch.Size([2, 2])
```

# TENSORS

Simply put, TENSORS are a generalization of vectors and matrices. In PyTorch, they are a multi-dimensional matrix containing elements of a **single data type**.

```
>>> import torch
>>> t = torch.tensor([[1., -1.], [1., -1.]])
>>> t
tensor([[ 1., -1.]
         [ 1., -1.]])

>>> t.dtype # They have a type
torch.float32

>>> t.shape # a shape
torch.Size([2, 2])

>>> t.device # and live in some device
device(type='cpu')
```

# TENSORS

- ▶ Although PyTorch has an elegant *python first* design, all PyTorch heavy work is actually implemented in C++.
- ▶ In Python, the integration of C++ code is (usually) done using what is called an **extension**;

# TENSORS

- ▶ Although PyTorch has an elegant *python first* design, all PyTorch heavy work is actually implemented in C++.
- ▶ In Python, the integration of C++ code is (usually) done using what is called an **extension**;
- ▶ PyTorch uses **ATen**, which is the foundational tensor operation library on which all else is built;

# TENSORS

- ▶ Although PyTorch has an elegant *python first* design, all PyTorch heavy work is actually implemented in C++.
- ▶ In Python, the integration of C++ code is (usually) done using what is called an **extension**;
- ▶ PyTorch uses **ATen**, which is the foundational tensor operation library on which all else is built;
- ▶ To do automatic differentiation, PyTorch uses **Autograd**, which is an augmentation on top of the **ATen** framework;

# TENSORS

- ▶ Although PyTorch has an elegant *python first* design, all PyTorch heavy work is actually implemented in C++.
- ▶ In Python, the integration of C++ code is (usually) done using what is called an **extension**;
- ▶ PyTorch uses **ATen**, which is the foundational tensor operation library on which all else is built;
- ▶ To do automatic differentiation, PyTorch uses **Autograd**, which is an augmentation on top of the **ATen** framework;
- ▶ In the Python API, PyTorch previously had separate **Variable** and a **Tensor** types, after v.0.4.0 they were merged into **Tensor**.



# QUICK RECAP PYTHON OBJECTS

```
typedef struct {  
    PyObject_HEAD  
    double ob_fval;  
} PyFloatObject;
```

# QUICK RECAP PYTHON OBJECTS

```
typedef struct {  
    PyObject_HEAD  
    double ob_fval;  
} PyFloatObject;
```

```
typedef struct _object {  
    Py_ssize_t ob_refcnt;  
    struct _typeobject *ob_type;  
} PyObject;
```

# QUICK RECAP PYTHON OBJECTS

```
typedef struct {
    PyObject_HEAD
    double ob_fval;
} PyFloatObject;
```


```
typedef struct _object {
    Py_ssize_t ob_refcnt;
    struct _typeobject *ob_type;
} PyObject;
```



PyFloatObject

object

PyObject\_HEAD  
double ob\_fval



PyObject

object

Py\_ssize\_t ob\_refcnt  
struct \_typeobject \*ob\_type

# QUICK RECAP PYTHON OBJECTS


```
struct THPVariable {
    PyObject_HEAD
    torch::autograd::Variable cdata;
    PyObject* backward_hooks;
};
```

---

The **TH** prefix is from **TorchH**, and **P** means **Python**.

# QUICK RECAP PYTHON OBJECTS

```
struct THPVariable {  
    PyObject_HEAD  
    torch::autograd::Variable cdata;  
    PyObject* backward_hooks;  
};
```



---

The **TH** prefix is from **TorchH**, and **P** means **Python**.

# IN PYTHON, EVERYTHING IS AN OBJECT

```
>>> a = 300
>>> b = 300
>>> a is b
False
```

# IN PYTHON, EVERYTHING IS AN OBJECT

```
>>> a = 300
```

```
>>> b = 300
```

```
>>> a is b
```

```
False
```

```
>>> a = 200
```

```
>>> b = 200
```

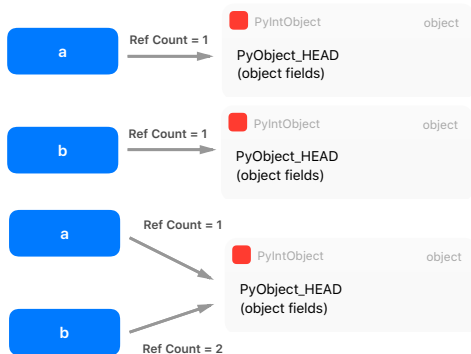
```
>>> a is b
```

```
True
```

# IN PYTHON, EVERYTHING IS AN OBJECT

```
>>> a = 300
>>> b = 300
>>> a is b
False

>>> a = 200
>>> b = 200
>>> a is b
True
```



A typical Python program spend much of its time allocating/deallocating integers. CPython then caches the small integers.



# ZERO-COPYING TENSORS

It is very common to load tensors in **numpy** and convert them to PyTorch, or vice-versa;

```
>>> np_array = np.ones((2,2))  
>>> np_array  
array([[1., 1.],  
       [1., 1.]])
```

---

Underline after an operation means an in-place operation.

# ZERO-COPYING TENSORS

It is very common to load tensors in **numpy** and convert them to PyTorch, or vice-versa;

```
>>> np_array = np.ones((2,2))
>>> np_array
array([[1., 1.],
       [1., 1.]])

>>> torch_array = torch.tensor(np_array)
>>> torch_array
tensor([[1., 1.],
        [1., 1.]], dtype=torch.float64)
```

---

Underline after an operation means an in-place operation.

# ZERO-COPYING TENSORS

It is very common to load tensors in **numpy** and convert them to PyTorch, or vice-versa;

```
>>> np_array = np.ones((2,2))
>>> np_array
array([[1., 1.],
       [1., 1.]])

>>> torch_array = torch.tensor(np_array)
>>> torch_array
tensor([[1., 1.],
        [1., 1.]], dtype=torch.float64)

>>> torch_array.add_(1.0)
```

---

Underline after an operation means an in-place operation.

# ZERO-COPYING TENSORS

It is very common to load tensors in **numpy** and convert them to PyTorch, or vice-versa;

```
>>> np_array = np.ones((2,2))
>>> np_array
array([[1., 1.],
       [1., 1.]])

>>> torch_array = torch.tensor(np_array)
>>> torch_array
tensor([[1., 1.],
        [1., 1.]], dtype=torch.float64)

>>> torch_array.add_(1.0)


>>> np_array
array([[1., 1.], # array is intact, a copy was made
       [1., 1.]])
```

---

Underline after an operation means an in-place operation.

# ZERO-COPYING TENSORS

- Now imagine that you have a batch of 128 images, 3 channels each (RGB) and with size of 224x224;



- This will yield a size in memory of ~ **74MB**. We don't want to duplicate memory (except when copying them to discrete GPUs of course);

# ZERO-COPYING TENSORS

Let's see now a slightly different code using the function `torch.from_numpy()` this time:

```
>>> np_array  
array([[1., 1.],  
       [1., 1.]])  
>>> torch_array = torch.from_numpy(np_array)
```

# ZERO-COPYING TENSORS

Let's see now a slightly different code using the function `torch.from_numpy()` this time:

```
>>> np_array
array([[1., 1.],
       [1., 1.]])
>>> torch_array = torch.from_numpy(np_array)
>>> torch_array.add_(1.0)
```

# ZERO-COPYING TENSORS

Let's see now a slightly different code using the function `torch.from_numpy()` this time:

```
>>> np_array
array([[1., 1.],
       [1., 1.]])
>>> torch_array = torch.from_numpy(np_array)
>>> torch_array.add_(1.0)
>>> np_array
array([[2., 2.],
       [2., 2.]])
```



# ZERO-COPYING TENSORS

Let's see now a slightly different code using the function `torch.from_numpy()` this time:

```
>>> np_array
array([[1., 1.],
       [1., 1.]])
>>> torch_array = torch.from_numpy(np_array)
>>> torch_array.add_(1.0)
>>> np_array
array([[2., 2.],
       [2., 2.]])
```

The original numpy array **was changed**, because it used a **zero-copy** operation.

# ZERO-COPYING TENSORS

Difference between **in-place** and **standard operations** might not be so clear in some cases:

```
>>> np_array  
array([[1., 1.],  
       [1., 1.]])  
>>> torch_array = torch.from_numpy(np_array)
```

# ZERO-COPYING TENSORS

Difference between **in-place** and **standard operations** might not be so clear in some cases:

```
>>> np_array  
array([[1., 1.],  
       [1., 1.]])  
>>> torch_array = torch.from_numpy(np_array)  
>>> np_array = np_array + 1.0
```

# ZERO-COPYING TENSORS

Difference between **in-place** and **standard operations** might not be so clear in some cases:

```
>>> np_array
array([[1., 1.],
       [1., 1.]])
>>> torch_array = torch.from_numpy(np_array)
>>> np_array = np_array + 1.0
>>> torch_array
tensor([[1., 1.],
        [1., 1.]], dtype=torch.float64)
```

# ZERO-COPYING TENSORS

Difference between **in-place** and **standard operations** might not be so clear in some cases:

```
>>> np_array
array([[1., 1.],
       [1., 1.]])
>>> torch_array = torch.from_numpy(np_array)
>>> np_array = np_array + 1.0
>>> torch_array
tensor([[1., 1.],
        [1., 1.]], dtype=torch.float64)
```

However, if you use `np_array += 1.0`, that is an in-place operation that will change `torch_array` memory.

# ZERO-COPYING TENSORS

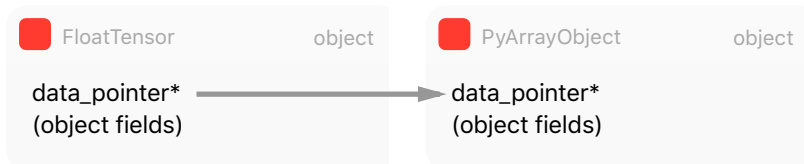
```

at::Tensor tensor_from_numpy(PyObject* obj) {
    // (...) - omitted for brevity
    auto array = (PyArrayObject*)obj;
    int ndim = PyArray_NDIM(array);
    auto sizes = to_aten_shape(ndim, PyArray_DIMS(array));
    auto strides = to_aten_shape(ndim, PyArray_STRIDES(array));
    // (...) - omitted for brevity
    void* data_ptr = PyArray_DATA(array);
    auto& type = CPU(dtype_to_aten(PyArray_TYPE(array)));
    Py_INCREF(obj);
    return type.tensorFromBlob(data_ptr, sizes, strides,
                                [obj](void* data) {
                                    AutoGIL gil;
                                    Py_DECREF(obj);
                                });
}

```

Pay attention to the reference counting using `Py_INCREF()` and the call to `tensorFromBlob()` function.

# DATA POINTERS



The tensor `FloatTensor` did a copy of the numpy array **data pointer** and not of the contents. The reference is kept safe by the Python reference counting mechanism.

# TENSOR STORAGE

The abstraction responsible for holding the data isn't actually the **Tensor**, but the **Storage**.



# TENSOR STORAGE

The abstraction responsible for holding the data isn't actually the **Tensor**, but the **Storage**.

```
struct C10_API StorageImpl final : (...) {  
  // (...)  
private:  
  // (...)  
  caffe2::TypeMeta data_type_;  
  DataPtr data_ptr_;  
  int64_t numel_;  
  Allocator* allocator_;  
}
```

# TENSOR STORAGE

The abstraction responsible for holding the data isn't actually the **Tensor**, but the **Storage**.

```
struct C10_API StorageImpl final : (...) {  
  // (...)  
private:  
  // (...)  
  caffe2::TypeMeta data_type_;  
  DataPtr data_ptr_;  
  int64_t numel_;  
  Allocator* allocator_;  
}
```

- ▶ Holds a pointer to the raw data and contains information such as the size and allocator;
- ▶ Storage is a dumb abstraction, there is no metadata telling us how to interpret the data it holds;

# TENSOR STORAGE

- The **Storage** abstraction is very powerful because it decouples the raw data and how we can interpret it;

# TENSOR STORAGE

- ▶ The **Storage** abstraction is very powerful because it decouples the raw data and how we can interpret it;
- ▶ We can have multiple tensors sharing the **same storage**, but with different interpretations, also called **views**, but **without duplicating** memory:

# TENSOR STORAGE

- ▶ The **Storage** abstraction is very powerful because it decouples the raw data and how we can interpret it;
- ▶ We can have multiple tensors sharing the **same storage**, but with different interpretations, also called **views**, but **without duplicating** memory:

```
>>> tensor_a = torch.ones((2, 2))
>>> tensor_b = tensor_a.view(4)
>>> tensor_a_data = tensor_a.storage().data_ptr()
>>> tensor_b_data = tensor_b.storage().data_ptr()
>>> tensor_a_data == tensor_b_data
True
```

# TENSOR STORAGE

- ▶ The **Storage** abstraction is very powerful because it decouples the raw data and how we can interpret it;
- ▶ We can have multiple tensors sharing the **same storage**, but with different interpretations, also called **views**, but **without duplicating** memory:

```
>>> tensor_a = torch.ones((2, 2))
>>> tensor_b = tensor_a.view(4)
>>> tensor_a_data = tensor_a.storage().data_ptr()
>>> tensor_b_data = tensor_b.storage().data_ptr()
>>> tensor_a_data == tensor_b_data
True
```

- ▶ **tensor\_b** is a different view (interpretation) of the same data present in the underlying storage that is shared between both tensors.

# MEMORY ALLOCATORS (CPU/GPU)

- The tensor storage can be allocated either in the CPU memory or GPU, therefore a mechanism is required to switch between these different allocations:

# MEMORY ALLOCATORS (CPU/GPU)

- ▶ The tensor storage can be allocated either in the CPU memory or GPU, therefore a mechanism is required to switch between these different allocations:

```
struct Allocator {
    virtual ~Allocator() {}
    virtual DataPtr allocate(size_t n) const = 0;
    virtual DeleterFnPtr raw_deleter() const {...}
    void* raw_allocate(size_t n) {...}
    void raw_deallocate(void* ptr) {...}
};
```

- ▶ There are `Allocator`s that will use the GPU allocators such as `cudaMallocHost()` when the storage should be used for the GPU or `posix_memalign()` POSIX functions for data in the CPU memory.



# THE BIG PICTURE



- The **Tensor** has a **Storage** which in turn has a pointer to the raw data and to the **Allocator** to allocate memory according to the destination device.

## Section II

∞ JIT ∞

# JIT - JUST-IN-TIME COMPILER

- ▶ PyTorch is eager by design, which means that it is easily hackable to debug, inspect, etc;

# JIT - JUST-IN-TIME COMPILER

- ▶ PyTorch is eager by design, which means that it is easily hackable to debug, inspect, etc;
- ▶ However, this poses problems for optimization and for decoupling it from Python (the model itself is Python code);

# JIT - JUST-IN-TIME COMPILER

- ▶ PyTorch is eager by design, which means that it is easily hackable to debug, inspect, etc;
- ▶ However, this poses problems for optimization and for decoupling it from Python (the model itself is Python code);
- ▶ PyTorch 1.0 introduced `torch.jit`, which has two main methods to convert a PyTorch model to a serializable and optimizable format;
- ▶ **TorchScript** was also introduced as a statically-typed subset of Python;

# JIT - JUST-IN-TIME COMPILER

Two very different worlds with their own requirements.

## ☐ EAGER MODE

Prototype, debug, train,  
experiment



tracing



scripting

## ☐ SCRIPT MODE

Optimization, other  
languages, deployment



# TRACING

```
def my_function(x):  
    if x.mean() > 1.0:  
        r = torch.tensor(1.0)  
    else:  
        r = torch.tensor(2.0)  
    return r
```

# TRACING

```
def my_function(x):  
    if x.mean() > 1.0:  
        r = torch.tensor(1.0)  
    else:  
        r = torch.tensor(2.0)  
    return r  
  
>>> ftrace = torch.jit.trace(my_function, (torch.ones(2, 2)))
```



# TRACING

```
def my_function(x):  
    if x.mean() > 1.0:  
        r = torch.tensor(1.0)  
    else:  
        r = torch.tensor(2.0)  
    return r  
  
>>> ftrace = torch.jit.trace(my_function, (torch.ones(2, 2)))  
  
>>> ftrace.graph  
graph(%x : Float(2, 2)) {  
  %4 : Float() = prim::Constant[value={2}]()  
  %5 : Device = prim::Constant[value="cpu"]()  
  %6 : int = prim::Constant[value=6]()  
  %7 : bool = prim::Constant[value=0]()  
  %8 : bool = prim::Constant[value=0]()  
  %9 : Float() = aten::to(%4, %5, %6, %7, %8)  
  %10 : Float() = aten::detach(%9)  
  return (%10); }
```

# TRACING

To call the JIT'ed function, just call the `forward()` method:

```
>>> x = torch.ones(2, 2)
>>> ftrace.forward(x)
tensor(2.)
```

# TRACING

To call the JIT'ed function, just call the `forward()` method:

```
>>> x = torch.ones(2, 2)
>>> ftrace.forward(x)
tensor(2.)
```

However, tracing will not record any control-flow like if statements or loops, it executes the code with the given context and creates the graph. You can see this limitation below:

```
>>> x = torch.ones(2, 2).add_(1.0)
>>> ftrace.forward(x)
tensor(2.)
```

According to `my_function()`, result should have been 1.0. Tracing also checks for differences between traced and Python function, but what about **Dropout**?

# SCRIPTING

Another alternative is to use **scripting**, where you can use decorators such as `@torch.jit.script`:

```
@torch.jit.script
def my_function(x):
    if bool(x.mean() > 1.0):
        r = 1
    else:
        r = 2
    return r
```

# SCRIPTING

```
>>> my_function.graph
graph(%x : Tensor) {
  %2 : float = prim::Constant[value=1]()
  %5 : int = prim::Constant[value=1]()
  %6 : int = prim::Constant[value=2]()
  %1 : Tensor = aten::mean(%x)
  %3 : Tensor = aten::gt(%1, %2)
  %4 : bool = prim::Bool(%3)
  %r : int = prim::If(%4)
    block0() {
      -> (%5)
    }
    block1() {
      -> (%6)
    }
  return (%r);
}
```

# SCRIPTING

The `my_function()` is now a `ScriptModule`:

```
>>> type(my_function)
torch.jit.ScriptModule
```

# SCRIPTING

The `my_function()` is now a `ScriptModule`:

```
>>> type(my_function)
torch.jit.ScriptModule
```

When we check the results again:

```
>>> x = torch.ones(2, 2)
>>> my_function(x)
2
```

# SCRIPTING

The `my_function()` is now a `ScriptModule` :

```
>>> type(my_function)
torch.jit.ScriptModule
```

When we check the results again:

```
>>> x = torch.ones(2, 2)
>>> my_function(x)
2
```

```
>>> x = torch.ones(2, 2).add_(1.0)
>>> my_function(x)
1
```

Control-flow logic was preserved !



# WHY TORCHSCRIPT ?

- The concept of having a well-defined **Intermediate Representation** (IR) is very powerful, it's the main concept behind LLVM platform as well;

# WHY TORCHSCRIPT ?

- ▶ The concept of having a well-defined **Intermediate Representation** (IR) is very powerful, it's the main concept behind LLVM platform as well;
- ▶ This opens the door to:
  - ▶ Decouple the model (computational graph) from Python runtime;

# WHY TORCHSCRIPT ?

- ▶ The concept of having a well-defined **Intermediate Representation** (IR) is very powerful, it's the main concept behind LLVM platform as well;
- ▶ This opens the door to:
  - ▶ Decouple the model (computational graph) from Python runtime;
  - ▶ Use it in production with C++ (no GIL) or other languages;

# WHY TORCHSCRIPT ?

- ▶ The concept of having a well-defined **Intermediate Representation** (IR) is very powerful, it's the main concept behind LLVM platform as well;
- ▶ This opens the door to:
  - ▶ Decouple the model (computational graph) from Python runtime;
  - ▶ Use it in production with C++ (no GIL) or other languages;
  - ▶ Capitalize on optimizations (whole program);

# WHY TORCHSCRIPT ?

- ▶ The concept of having a well-defined **Intermediate Representation** (IR) is very powerful, it's the main concept behind LLVM platform as well;
- ▶ This opens the door to:
  - ▶ Decouple the model (computational graph) from Python runtime;
  - ▶ Use it in production with C++ (no GIL) or other languages;
  - ▶ Capitalize on optimizations (whole program);
  - ▶ Split the development world of hackable and easy to debug from the world of putting these models in production and optimize them.

# BUILDING THE IR


To build the IR, PyTorch takes leverage of the Python **Abstract Syntax Tree** (AST) which is a tree representation of the syntactic structure of the source code.

```
>>> ast_mod = ast.parse("print(1 + 2)")
>>> astpretty.pprint(ast_mod.body[0], show_offsets=False)
```

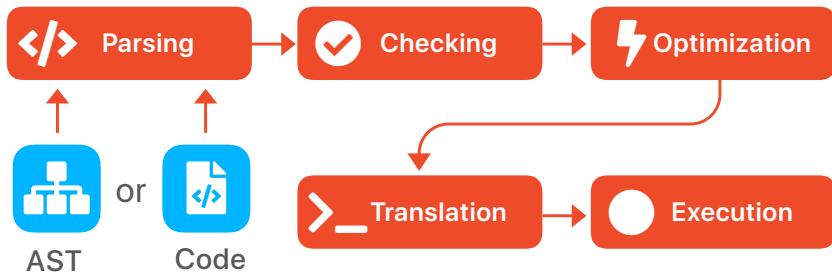
```
Expr(
  value=Call(
    func=Name(id='print', ctx=Load()),
    args=[
      BinOp(
        left=Num(n=1),
        op=Add(),
        right=Num(n=2),
      ),
    ],
    keywords=[],
  ),
)
```

# BUILDING THE IR

```
print(1 + 2)
```



# PYTORCH JIT PHASES





# EXECUTING

Just like Python interpreter executes your code, PyTorch has a interpreter that executes the IR instructions:

```
bool runImpl(Stack& stack) {
    auto& instructions = function->instructions;
    size_t last = instructions.size();

    while (pc < last) {
        auto& inst = instructions[pc];
        try {
            loadTensorsFromRegisters(inst.inputs, stack);
            size_t new_pc = pc + 1 + inst.callback(stack);
            for (int i = inst.outputs.size - 1; i >= 0; --i) {
                int reg = get(inst.outputs, i);
                registers[reg] = pop(stack);
            }
            pc = new_pc;

            // (...) omitted
        }
    }
}
```

# OPTIMIZATIONS

Many optimizations can be used on the computational graph of the model, such as **Loop Unrolling**:

```
for i.. i+= 1
  for j..
    code(i, j)
```

```
for i.. i+= 4
  for j..
    code(i, j)
    code(i+1, j)
    code(i+2, j)
    code(i+3, j)
  remainder loop
```

# OPTIMIZATIONS

Also **Peephole optimizations** such as:

$$x.t().t() = x$$

# OPTIMIZATIONS

Also **Peephole optimizations** such as:

```
x.t().t() = x
```

Example:

```
def dumb_function(x):  
    return x.t().t()
```

```
>>> traced_fn = torch.jit.trace(dumb_function,  
...                             torch.ones(2,2))  
>>> traced_fn.graph_for(torch.ones(2,2))  
graph(%x : Float(*, *)) {  
    return (%x);  
}
```

# OPTIMIZATIONS

Also **Peephole optimizations** such as:

$$x.t().t() = x$$

Example:

```
def dumb_function(x):
    return x.t().t()

>>> traced_fn = torch.jit.trace(dumb_function,
...                             torch.ones(2,2))
>>> traced_fn.graph_for(torch.ones(2,2))
graph(%x : Float(*, *)) {
  return (%x);
}
```

Other optimizations include **Constant Propagation**, **Dead Code Elimination (DCE)**, **fusion**, **inlining**, etc.

# SERIALIZATION

```
>>> resnet = torch.jit.trace(models.resnet18(),  
...                           torch.rand(1, 3, 224, 224))  
>>> resnet.save("resnet.pt")
```

# SERIALIZATION

```
>>> resnet = torch.jit.trace(models.resnet18(),  
...                           torch.rand(1, 3, 224, 224))  
>>> resnet.save("resnet.pt")  
  
$ file resnet.pt  
resnet.pt: Zip archive data
```

# SERIALIZATION

```
>>> resnet = torch.jit.trace(models.resnet18(),  
...                           torch.rand(1, 3, 224, 224))  
>>> resnet.save("resnet.pt")
```

```
$ file resnet.pt
```

```
resnet.pt: Zip archive data
```

```
$ unzip resnet.pt
```

```
Archive:  resnet.pt
```

```
extracting: resnet/version
```

```
extracting: resnet/code/resnet.py
```

```
extracting: resnet/model.json
```

```
extracting: resnet/tensors/0
```

```
(...)
```



# SERIALIZATION

code/resnet.py

```
op_version_set = 0
```

```
def forward(self, input_1: Tensor) -> Tensor:
```

```
    input_2 = torch._convolution(input_1, self.conv1.weight, ...)
```

```
    # (...)
```

```
    input_3 = torch.batch_norm(input_2, self.bn1.weight, self.bn1.bias,  
                               self.bn1.running_mean, self.bn1.running_var, ...)
```

```
    # (...)
```

# SERIALIZATION

code/resnet.py

```
op_version_set = 0
def forward(self, input_1: Tensor) -> Tensor:
    input_2 = torch._convolution(input_1, self.conv1.weight, ...)
    # (...)
    input_3 = torch.batch_norm(input_2, self.bn1.weight, self.bn1.bias,
                               self.bn1.running_mean, self.bn1.running_var, ...)
    # (...)
```

model.json

```
{"parameters":
  [{ "isBuffer": false,
    "tensorId": "1",
    "name": "weight" }],
  "name": "conv1",
  "optimize": true}
```

# SERIALIZATION

code/resnet.py

```
op_version_set = 0
def forward(self, input_1: Tensor) -> Tensor:
    input_2 = torch._convolution(input_1, self.conv1.weight, ...)
    # (...)
    input_3 = torch.batch_norm(input_2, self.bn1.weight, self.bn1.bias,
                               self.bn1.running_mean, self.bn1.running_var, ...)
    # (...)
```

model.json

```
{"parameters":
  [{ "isBuffer": false,
    "tensorId": "1",
    "name": "weight" }],
  "name": "conv1",
  "optimize": true}
```

model.json

```
[{"isBuffer": true,
  "tensorId": "4",
  "name": "running_mean"},
 {"isBuffer": true,
  "tensorId": "5",
  "name": "running_var"}],
  "name": "bn1",
  "optimize": true}
```

## USING THE MODEL IN C++

PyTorch also has a C++ API that you can use to load/train models in C++. This is good for production, mobile, embedded devices, etc.

Example of loading a traced model in PyTorch C++ API:

```
#include <torch/script.h>
int main(int argc, const char* argv[])
{
    auto module = torch::jit::load("resnet.pt");
    std::vector<torch::jit::IValue> inputs;
    inputs.push_back(torch::ones({1, 3, 224, 224}));
    at::Tensor output = module->forward(inputs).toTensor();
}
```

# USING THE MODEL IN NODEJS



```
> var torchjs = require("torchjs");
> var script_module = new torchjs.ScriptModule("resnet18_trace.pt");
> var data = torchjs.ones([1, 3, 224, 224], false);
> console.log(data);
Tensor[Type=Variable[CPUFloatType], Size=[1, 3, 224, 224]

> var output = script_module.forward(data);
> console.log(output);
Tensor[Type=Variable[CPUFloatType], Size=[1, 1000]
```

Complete tutorial at <https://goo.gl/7wMJuS>.

## Section III

# ∞ PRODUCTION ∞

# ISSUES WITH TUTORIALS

- ▶ Be careful with online tutorials using Flask, etc. They are simple, but they often fail on good practices:
  - ▶ They often use JSON and base64 to serialize images. This adds ~ 33% overhead **per call** (uncompressed);

# ISSUES WITH TUTORIALS

- ▶ Be careful with online tutorials using Flask, etc. They are simple, but they often fail on good practices:
  - ▶ They often use JSON and base64 to serialize images. This adds  $\sim 33\%$  overhead **per call** (uncompressed);
  - ▶ They don't pay attention to zero-copy practices, so they often transform, reshape, convert to numpy, convert to PyTorch, etc;



# ISSUES WITH TUTORIALS

- ▶ Be careful with online tutorials using Flask, etc. They are simple, but they often fail on good practices:
  - ▶ They often use JSON and base64 to serialize images. This adds  $\sim 33\%$  overhead **per call** (uncompressed);
  - ▶ They don't pay attention to zero-copy practices, so they often transform, reshape, convert to numpy, convert to PyTorch, etc;
  - ▶ They often use HTTP/1;

# ISSUES WITH TUTORIALS

- ▶ Be careful with online tutorials using Flask, etc. They are simple, but they often fail on good practices:
  - ▶ They often use JSON and base64 to serialize images. This adds  $\sim 33\%$  overhead **per call** (uncompressed);
  - ▶ They don't pay attention to zero-copy practices, so they often transform, reshape, convert to numpy, convert to PyTorch, etc;
  - ▶ They often use HTTP/1;
  - ▶ They seldom do batching (important for GPUs);

# ISSUES WITH TUTORIALS

- ▶ Be careful with online tutorials using Flask, etc. They are simple, but they often fail on good practices:
  - ▶ They often use JSON and base64 to serialize images. This adds  $\sim 33\%$  overhead **per call** (uncompressed);
  - ▶ They don't pay attention to zero-copy practices, so they often transform, reshape, convert to numpy, convert to PyTorch, etc;
  - ▶ They often use HTTP/1;
  - ▶ They seldom do batching (important for GPUs);
  - ▶ They never put that "production" code in production.

# PREFER BINARY SERIALIZATION FORMATS

Prefer using good **binary serialization** methods such as Protobuf that offers a **schema** and a schema evolution mechanism.

Example from EuclidesDB RPC message:

```
message AddImageRequest {  
    int32 image_id = 1;  
    bytes image_data = 2;  
    // This field can encode JSON data  
    bytes image_metadata = 3;  
    repeated string models = 4;  
}
```

---

\* <http://euclidesdb.readthedocs.io>

# AVOID EXTRA COPIES

- Be careful to avoid extra copies of your tensors, especially during pre-processing;

# AVOID EXTRA COPIES

- ▶ Be careful to avoid extra copies of your tensors, especially during pre-processing;
- ▶ You can use in-place operations. It is a functional anti-pattern because it introduces side-effects, but it's a fair price to pay for performance;

# AVOID EXTRA COPIES

- ▶ Be careful to avoid extra copies of your tensors, especially during pre-processing;
- ▶ You can use in-place operations. It is a functional anti-pattern because it introduces side-effects, but it's a fair price to pay for performance;
- ▶ **Caveat:** in-place operations doesn't make much sense when you need gradients. PyTorch uses tensor versioning to catch that:

```
>>> a = torch.tensor(1.0, requires_grad=True)
```

```
>>> y = a.tanh()
```

```
>>> y.add_(2.0)
```

```
>>> y.backward() # error !
```


```
>>> a._version
```

```
0
```

```
>>> y._version
```


```
1
```

# A TALE OF TWO HTTPs







# A TALE OF TWO HTTPs




# A TALE OF TWO HTTPs



# A TALE OF TWO HTTPs




# A TALE OF TWO HTTPs



- Use HTTP 2.0 if possible, and avoid the *head-of-line blocking*;


# A TALE OF TWO HTTPs



- ▶ Use HTTP 2.0 if possible, and avoid the *head-of-line blocking*;
- ▶ Even better, you can use frameworks such as gRPC that uses HTTP/2.0 and Protobuf.

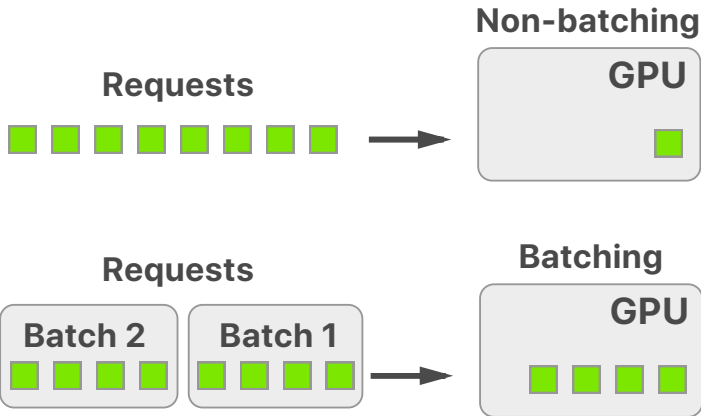
# BATCHING

Batching data is a way to amortize the performance bottleneck.



# BATCHING

Batching data is a way to amortize the performance bottleneck.



## Section IV

∞ Q&A ∞



# Q&A

Thanks !

